Labs
Optimization for Machine Learning
Spring 2021

EPFL
School of Computer and Communication Sciences
Martin Jaggi & Nicolas Flammarion
github.com/epfml/OptML_course

Problem Set 9 — Solutions (Coordinate Descent)

Exercise 56. *Let f be smooth with constant L in the classical sense, and satisfy the PL inequality (10.4). Let the problem $\min_{\mathbf{x}} f(\mathbf{x})$ have a non-empty solution set $\mathcal{X}^\star$. Prove that* gradient descent *with a stepsize of $1/L$ has a global linear convergence rate*

$$f(\mathbf{x}_t) - f^\star \leq \Big(1 - \frac{\mu}{L}\Big)^t (f(\mathbf{x}_0) - f^\star).$$

Solution: We combine smoothness $f(\mathbf{y}) \leq f(\mathbf{x}) + \langle \nabla f(\mathbf{x}), \mathbf{y} - \mathbf{x}\rangle + \frac{L}{2}\|\mathbf{y} - \mathbf{x}\|^2$ with the gradient descent update $\mathbf{x}_{t+1} := \mathbf{x}_t - \frac{1}{L}\nabla f(\mathbf{x}_t)$.

We then obtain

$$\begin{aligned} f(\mathbf{x}_{t+1}) &\leq f(\mathbf{x}_t) + \langle \nabla f(\mathbf{x}_t), \mathbf{x}_{t+1} - \mathbf{x}_t\rangle + \frac{L}{2}\|\mathbf{x}_{t+1} - \mathbf{x}_t\|^2 \\ &= f(\mathbf{x}_t) - \frac{1}{2L}\|\nabla f(\mathbf{x}_t)\|^2 \\ &\leq f(\mathbf{x}_t) - \frac{\mu}{L}(f(\mathbf{x}_t) - f^\star) \end{aligned}$$

the last two lines being the PL inequality. If we subtract $f^\star$ and apply recursively, we obtain the claimed linear rate:

$$f(\mathbf{x}_t) - f^\star \leq \Big(1 - \frac{\mu}{L}\Big)^t (f(\mathbf{x}_0) - f^\star).$$

Exercise 57. *Consider random coordinate descent with selecting the i-th coordinate with probability proportional to the L_i value, where L_i is the individual smoothness constant for each coordinate i as in* (10.5).

When using a stepsize of $1/L_{i_t}$, prove that we obtain the faster rate of

$$\mathbb{E}[f(\mathbf{x}_t) - f^\star] \leq \Big(1 - \frac{\mu}{d\bar{L}}\Big)^t [f(\mathbf{x}_0) - f^\star],$$

where $\bar{L} = \frac{1}{d}\sum_{i=1}^d L_i$ now is the average of all coordinate-wise smoothness constants. Note that this value can be much smaller than the global L we have used above, since that one was required to hold for all i so has to be chosen as $L = \max_i L_i$ instead.

Can you come up with an example from machine learning where $\bar{L} \ll L$?

Solution: We combine the coordinate descent update $\mathbf{x}_{t+1} := \mathbf{x}_t - \frac{1}{L_{i_t}}\nabla_{i_t} f(\mathbf{x}_t)\mathbf{e}_{i_t}$, with the property of coordinate-wise smoothness, $f(\mathbf{x} + \gamma\mathbf{e}_i) \leq f(\mathbf{x}) + \gamma\nabla_i f(\mathbf{x}) + \frac{L_i}{2}\gamma^2$, resulting in

$$\begin{aligned} f(\mathbf{x}_{t+1}) &= f(\mathbf{x}_t - \frac{1}{L_{i_t}}\nabla_{i_t} f(\mathbf{x}_t)\mathbf{e}_{i_t}) \\ &\leq f(\mathbf{x}) - \frac{1}{L_{i_t}}|\nabla_{i_t} f(\mathbf{x})|^2 + \frac{1}{2L_i}|\nabla_{i_t} f(\mathbf{x})|^2 \\ &= f(\mathbf{x}) - \frac{1}{2L_{i_t}}|\nabla_{i_t} f(\mathbf{x})|^2. \end{aligned}$$

Taking the average we have:

$$\begin{aligned} \mathbb{E}[f(\mathbf{x}_{t+1})] &\leq f(\mathbf{x}_t) - \frac{1}{2}\mathbb{E}\Big[\frac{1}{L_{i_t}}|\nabla_{i_t} f(\mathbf{x}_t)|^2\Big] \\ &= f(\mathbf{x}_t) - \frac{1}{2}\sum_i \Big[\frac{L_i}{\sum_j L_j}\frac{1}{L_i}|\nabla_i f(\mathbf{x}_t)|^2\Big] \\ &= f(\mathbf{x}_t) - \frac{1}{2\sum_j L_j}\|\nabla f(\mathbf{x}_t)\|^2 \end{aligned}$$

Exercise 58. *Derive the solution to exact coordinate minimization for the Lasso problem* (10.8)*, for the i-th coordinate. Write A_{-i} for the $n \times (d-1)$ matrix obtained by removing the i-th column from A, and same for the vector $\mathbf{x}_{-i}$ with one entry removed accordingly.*

Solution: We use the subgradient optimality condition for unconstrained convex minimization, applied to the single coordinate problem. A subgradient of this univariate objective can be written as

$$\begin{aligned}\frac{\partial}{\partial x_i}\left[\|A\mathbf{x}-\mathbf{b}\|^2+\lambda\|\mathbf{x}\|_1\right] &= 2A_i^\top[A\mathbf{x}-\mathbf{b}]+\lambda s \\ &= 2A_i^\top A_i x_i + 2A_i^\top(A_{-i}\mathbf{x}_{-i}-\mathbf{b})+\lambda s\end{aligned}$$

for $s \in \partial|x_i|$ being any subgradient of the (again univariate) absolute value function, and A_i is the i^{th} column of A.

At optimality, the previous partial derivative equals to zero, i.e. $0 \overset{!}{=} 2A_i^\top A_i x_i + 2A_i^\top(A_{-i}\mathbf{x}_{-i}-\mathbf{b})+\lambda s$.

Solving for x_i, this gives us:

$$\begin{aligned}x_i &= \frac{-A_i^\top(A_{-i}\mathbf{x}_{-i}-\mathbf{b})-\frac{1}{2}\lambda s}{A_i^\top A_i} \\ &= \frac{A_i^\top(\mathbf{b}-A_{-i}\mathbf{x}_{-i})}{\|A_i\|^2}-\frac{\lambda s}{2\|A_i\|^2} \\ &= S_{\frac{\lambda/2}{\|A_i\|^2}}\Big(\frac{A_i^\top(\mathbf{b}-A_{-i}\mathbf{x}_{-i})}{\|A_i\|^2}\Big)\end{aligned}$$

The left S operator corresponds to soft thresholding, defined as

$$S_a(b) := \begin{cases}0, & |b| \le a, \\ b-a & b > a, \\ b+a & b < -a\end{cases}.$$